

# Deployment with Docker and AWS

Updated Documentation Spring 2024

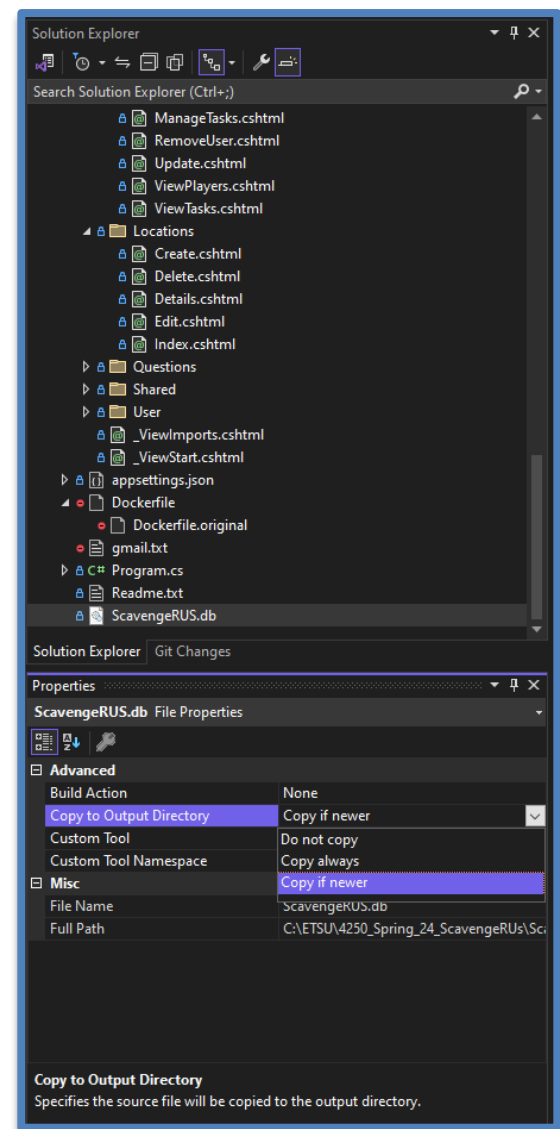
## Prerequisites

- Create a Docker account at <https://hub.docker.com>
- Download Docker Desktop <https://www.docker.com/products/docker-desktop/>
  - You will have to log in
- Create an AWS account at <https://portal.aws.amazon.com/billing/signup>
  - You will have to provide credit card info. Free tier provides an allowance of time before charging. Make sure to stop your instance and this should not be a problem.
- (Optional) This is for Windows machines, steps will differ for Mac/Linux

## Prep Database for Docker (For Development)

This technique will ensure that the database file already in the project files will retain its data when published to Docker as an image. This is useful to Developers as it can be used to quickly test database transactions with pre-populated data. However, there is a caveat: each time the image is deployed to a container, it will only contain the data from when the image was originally published

1. In the solution explorer, click the ScavengeRUS.db (database) file
2. Look at the option Copy to Output Directory and set it to *Copy if newer* or *Copy always*



## Visual Studio to Docker

1. In Visual Studio, while in your project navigate to the top bar and click “Build” and then “Publish <Project name>”
  - If you do not see a publish with Docker option, make sure you open the project by double-clicking the `.sln` or `.csproj` file (opens project as a solution)

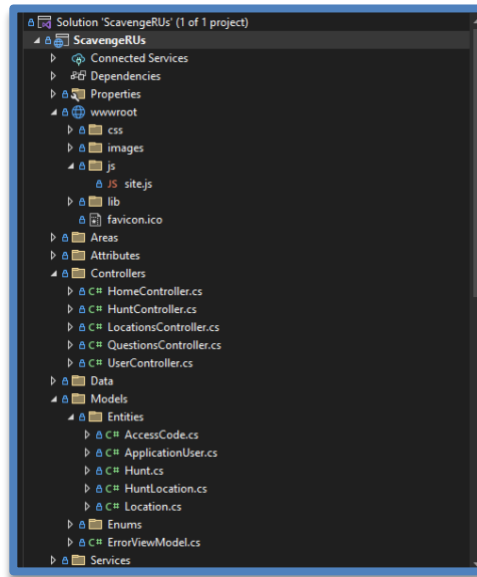
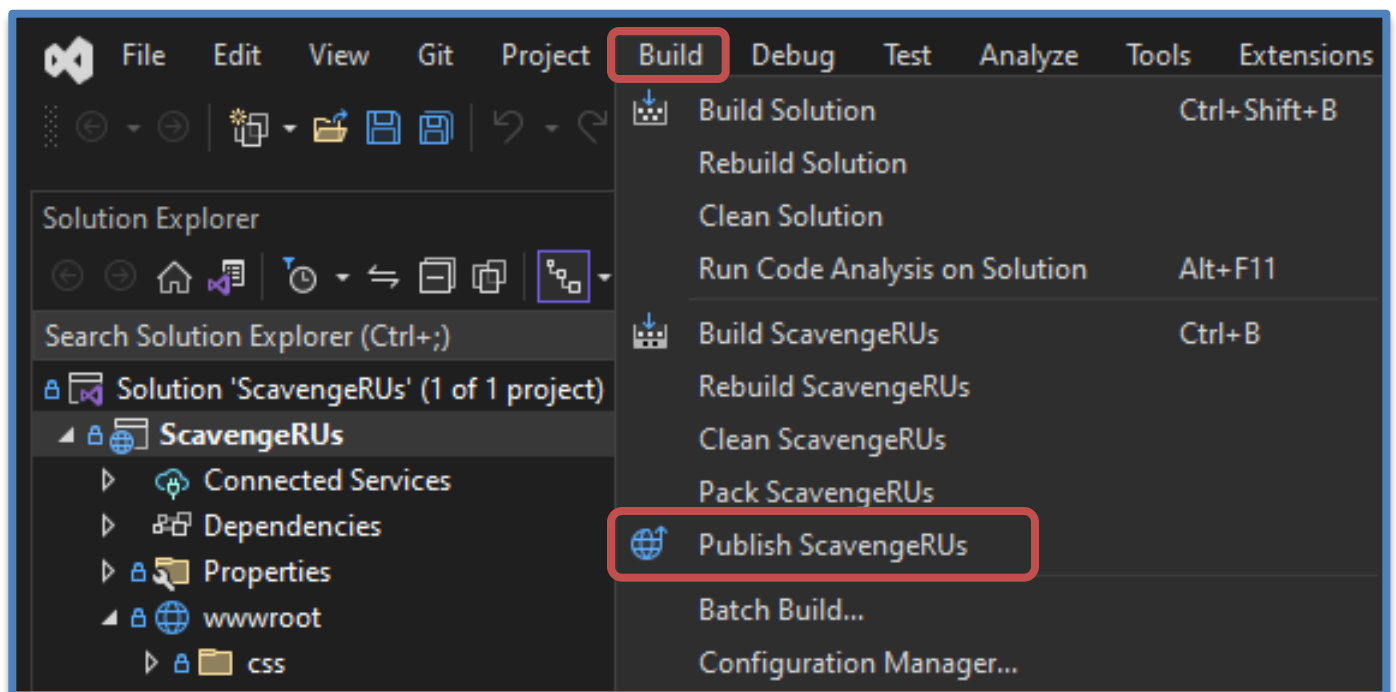
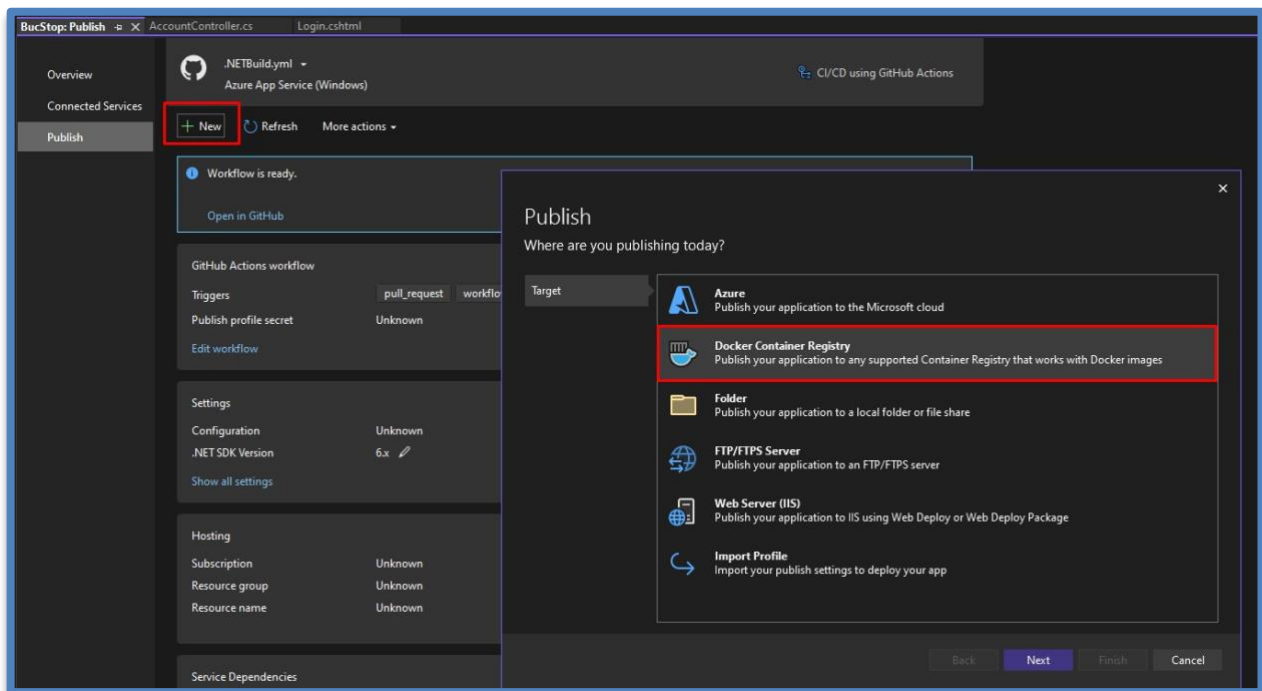


Figure 1 - You should see something like this if you opened it correctly

- If you are on Mac/Linux, Visual Studio support is sparse or no-longer supported. You can right click to add Docker Support, but you will not be able to deploy to AWS or to your Docker account

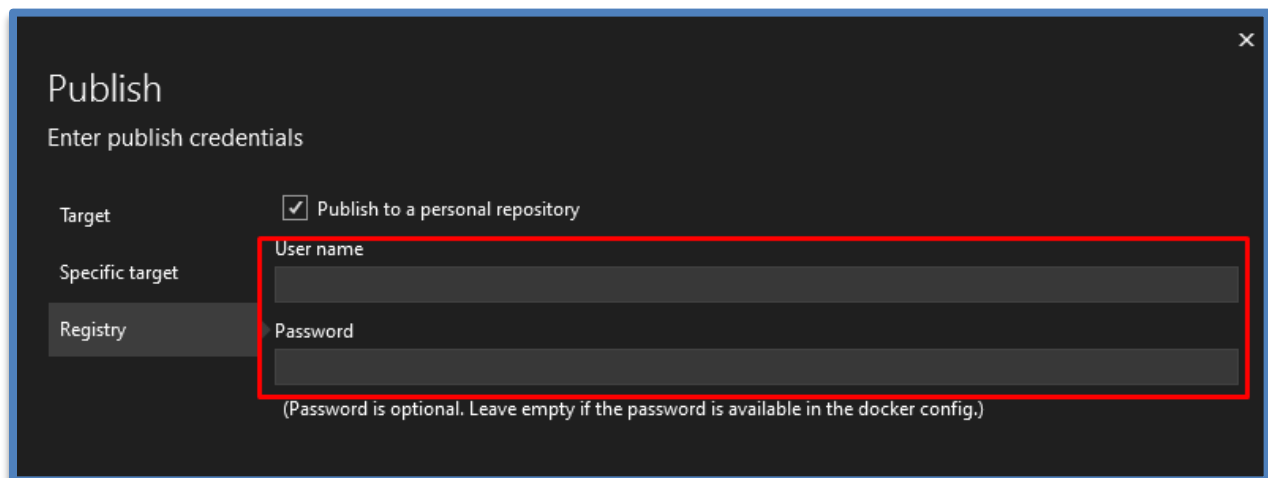


2. If this is the first time, you have to press **New** and then choose your target on where you want to publish, in this case it will be **Docker**. Once you press next, you will have three options to choose from, choose the one that says **Docker Hub**

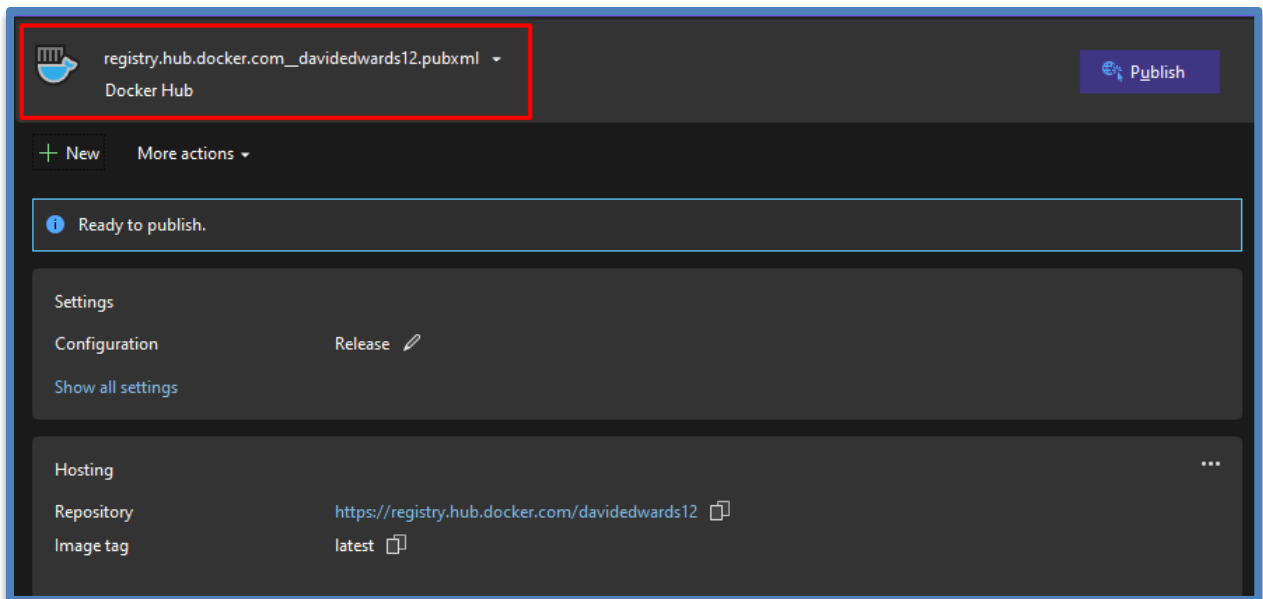


Note: You might have to install some packages in your VS if you do not have the required packages. If this is the case, go ahead and install them and then continue.

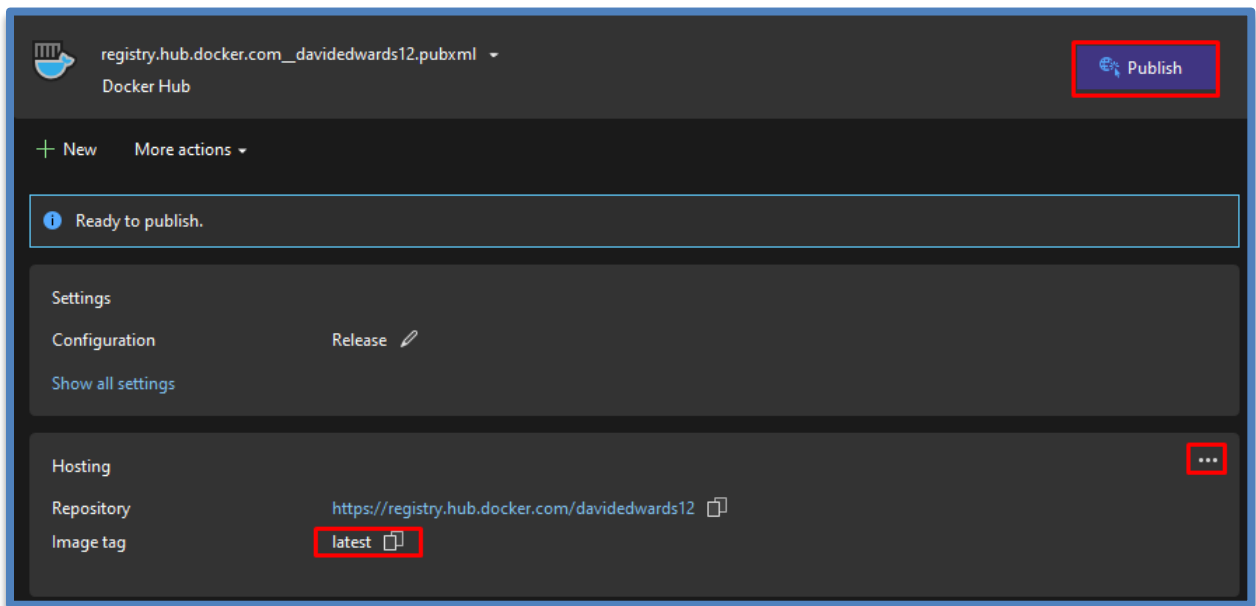
3. Next, log in into your Docker Hub account that you created at the start. Once you log in, you can press **Close**



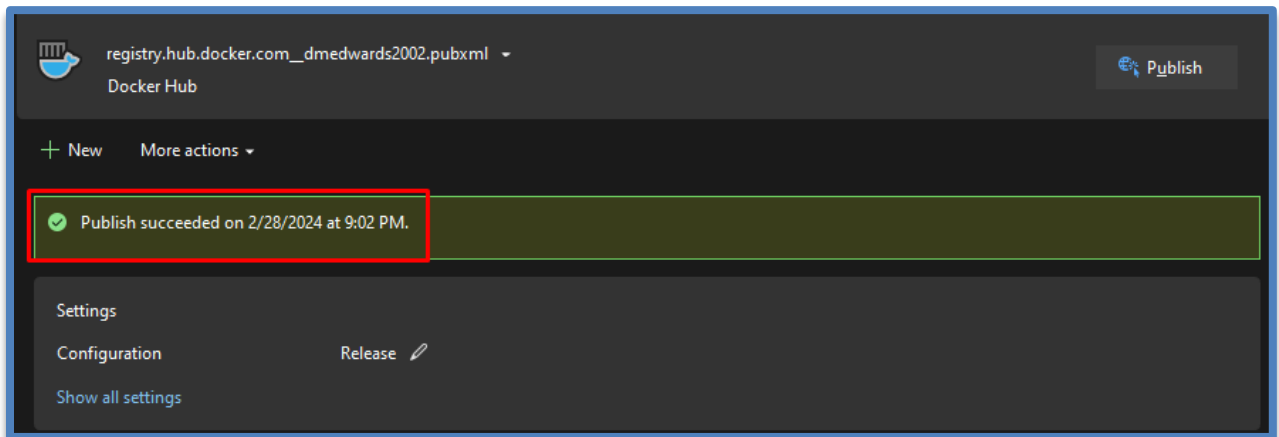
- Back in the publish tab, you should see that you are now logged into your Docker Hub account and can publish to your Docker Hub



- Now we can publish the project! Here, you can change the image tag by pressing the three little dots to the side (I think of the image tag as the version of the project, i.e., 1.0.0, 1.0.1, etc..). You can leave it latest if you want to, just know that it replaces the last latest you had. Once you have the image tag set to whatever you want, you can now press the “Publish button” at the top (if you have not downloaded Docker Desktop by this point, you will need to now as VS requires it to publish)

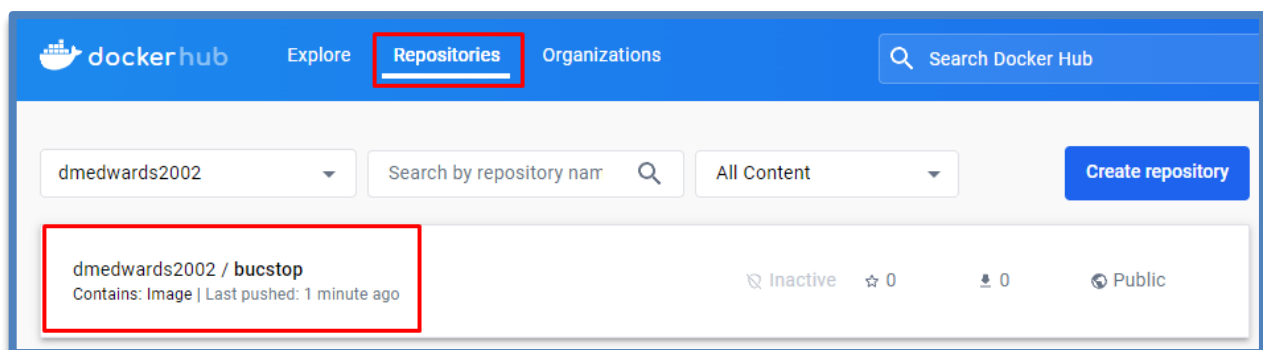
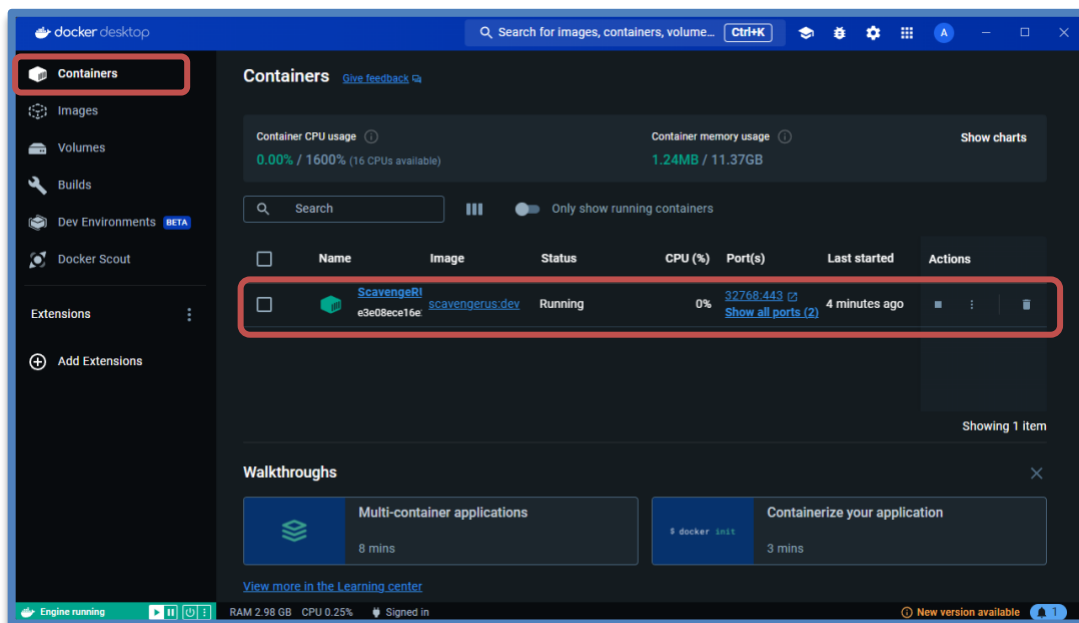


6. You will know you have published if you see this



7. Once you have that, you can go to your Docker Hub account and under repositories you should see your newly published project. You should be able to see this in the Docker Desktop as well

- If you do not have both of these, try running it again from the log-in step onward

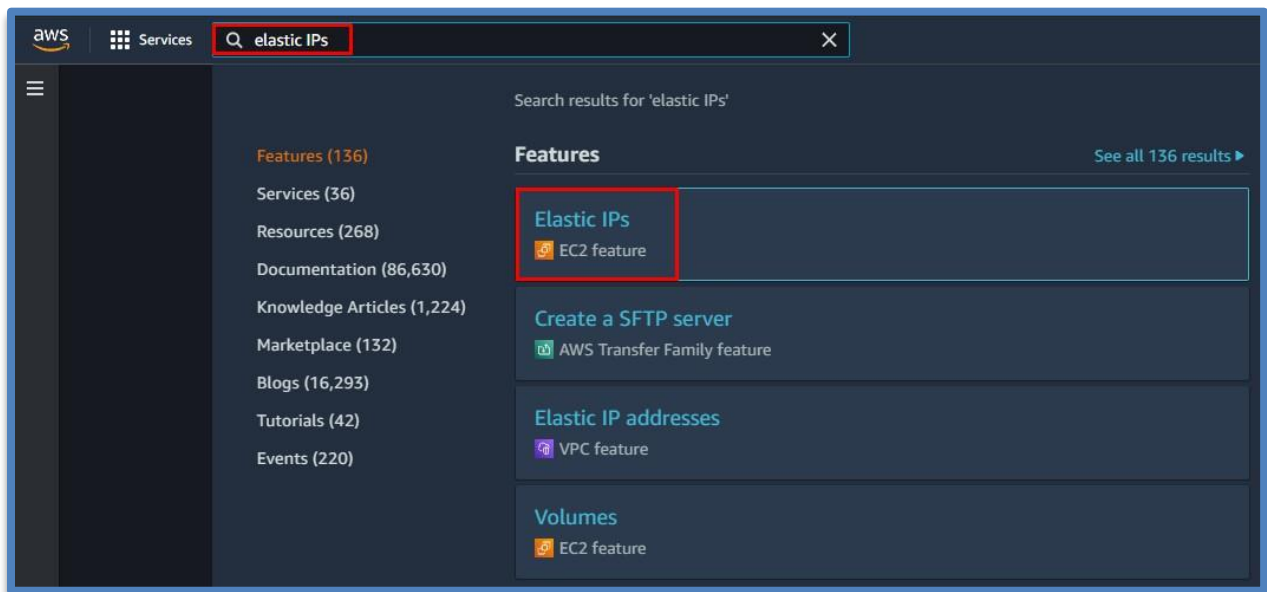


You can now deploy two different ways

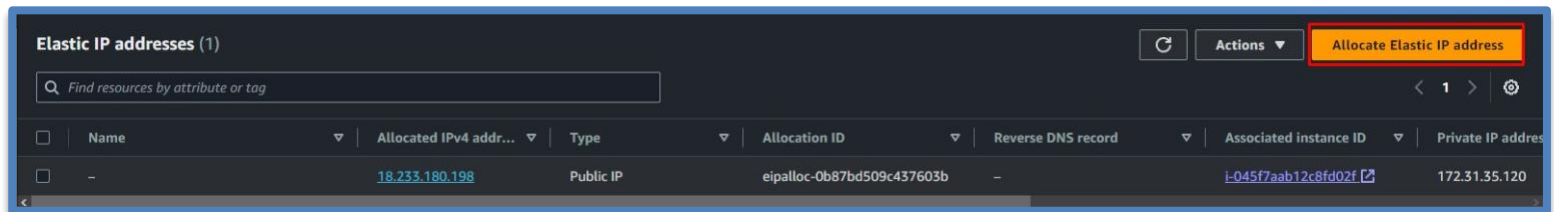
- Locally through the Docker Desktop application that we downloaded earlier
- Go through a public IP address with AWS. For this demo, we will be using AWS and an elastic IP address (more on that later).

## Setting Up AWS

1. Assuming you have already made an AWS account, head over to the search bar at the top and search for “Elastic IPs”, we will use this later.



2. In the top right, click on **Allocate Elastic IP address**



3. I recommend leaving everything as it is and pressing **Allocate** at the button. You will be given an IP address to use when we set up our virtual server

aws Services Search [Alt+S]

### Elastic IP address settings [Info](#)

Network border group [Info](#)

Q us-east-1 X

Public IPv4 address pool

- ☒ Amazon's pool of IPv4 addresses
- ☐ Public IPv4 address that you bring to your AWS account with BYOIP. (option disabled because no pools found) [Learn more](#)
- ☐ Customer-owned pool of IPv4 addresses created from your on-premises network for use with an Outpost. (option disabled because no customer owned pools found) [Learn more](#)

Global static IP addresses

AWS Global Accelerator can provide global static IP addresses that are announced worldwide using anycast from AWS edge locations. This can help improve the availability and latency for your user traffic by using the Amazon global network. [Learn more](#)

Create accelerator

### Tags - optional

A tag is a label that you assign to an AWS resource. Each tag consists of a key and an optional value. You can use tags to search and filter your resources or track your AWS costs.

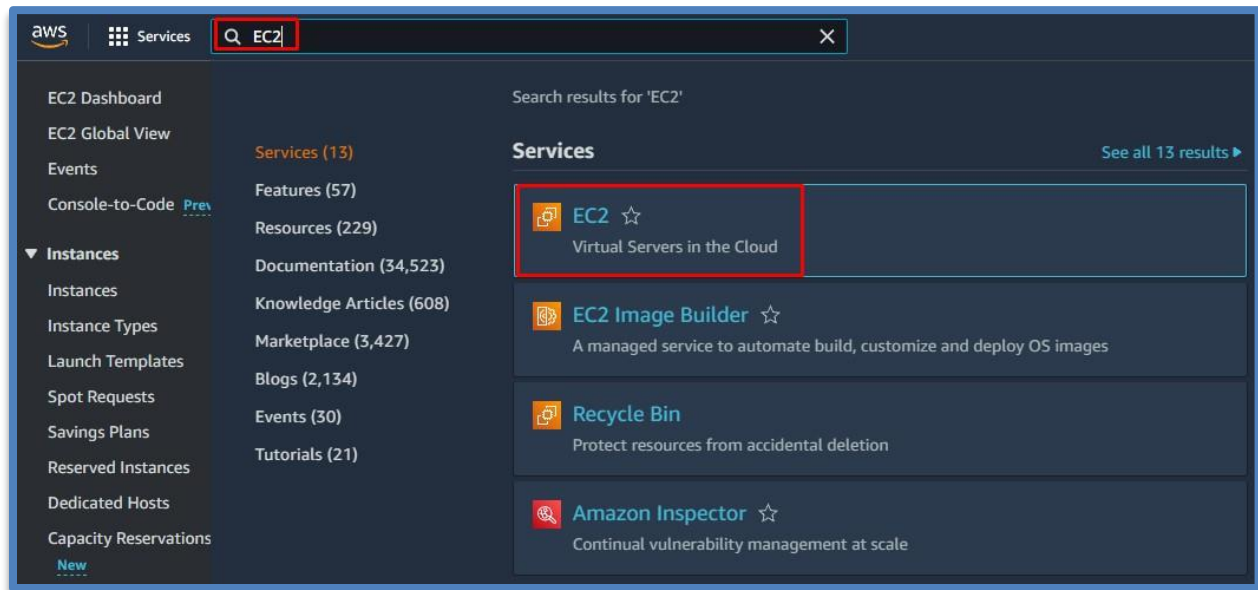
No tags associated with the resource.

Add new tag

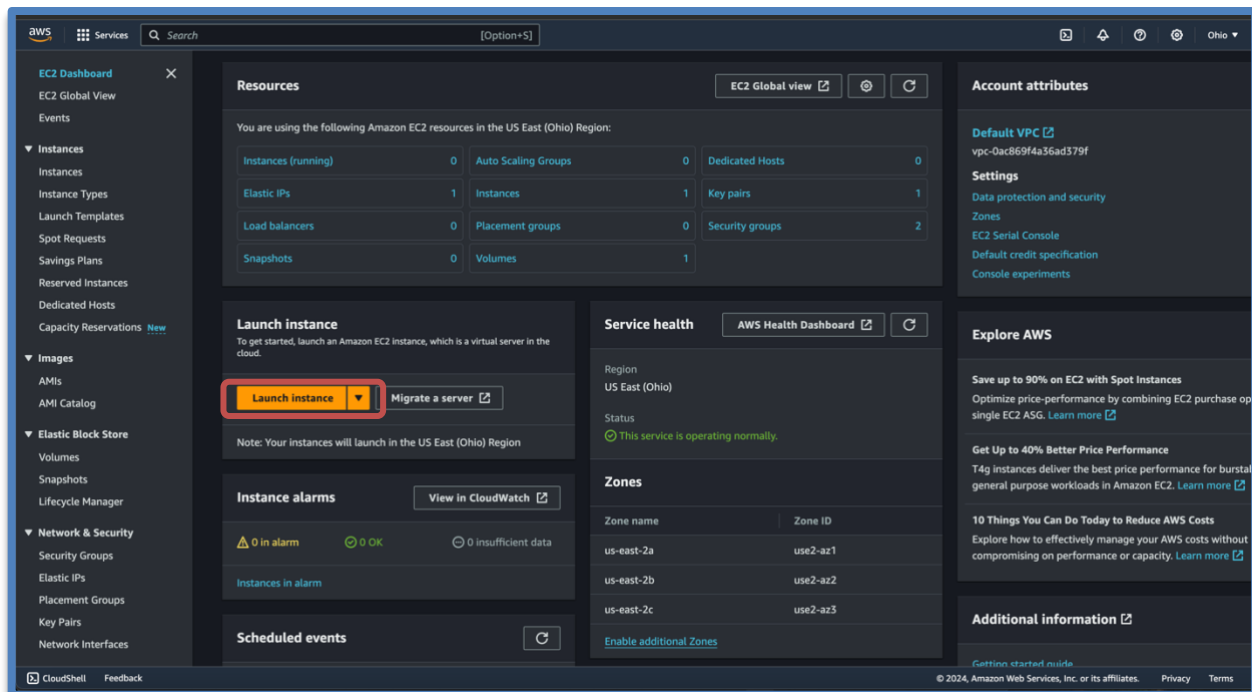
You can add up to 50 more tag

Cancel Allocate

4. Next, we search for “EC2.” This is where we set up our virtual server



5. On that page, you will want to press **Launch instance**





- Now we are setting up the instance. Name it and define an image. For this we will be using Amazon Linux

### Name and tags Info

Name

Test

Add additional tags

### ▼ Application and OS Images (Amazon Machine Image) Info

An AMI is a template that contains the software configuration (operating system, application server, and applications) required to launch your instance. Search or Browse for AMIs if you don't see what you are looking for below

Q Search our full catalog including 1000s of application and OS images

Recents

Quick Start

Amazon Linux  
aws

macOS  
Mac

Ubuntu  
ubuntu

Windows  
Microsoft

Red Hat  
RedHat

SUSE Li  
SUSE

Q

Browse more AMIs

Including AMIs from AWS, Marketplace and the Community

#### Amazon Machine Image (AMI)

Amazon Linux 2023 AMI  
ami-0440d3b780d96b29d (64-bit (x86), uefi-preferred) / ami-0f93c02efd1974b8b (64-bit (Arm), uefi)  
Virtualization: hvm   ENA enabled: true   Root device type: ebs

Free tier eligible ▼

- (Optional) Scrolling down, you can use a Key Pair to secure the instance more. It is not necessary but recommended. Then, in the network settings, check all the boxes. This will produce a certificate to download and hold on to.

## ▼ Key pair (login) [Info](#)

You can use a key pair to securely connect to your instance. Ensure that you have access to the selected key pair before you launch the instance.

Key pair name - *required*

etsu2024



[Create new key pair](#)

## ▼ Network settings [Info](#)

[Edit](#)

Network [Info](#)

vpc-063dc5898d57461d0

Subnet [Info](#)

No preference (Default subnet in any availability zone)

Auto-assign public IP [Info](#)

Enable

Firewall (security groups) [Info](#)

A security group is a set of firewall rules that control the traffic for your instance. Add rules to allow specific traffic to reach your instance.

☒ Create security group

☐ Select existing security group

We'll create a new security group called 'launch-wizard-8' with the following rules:

☒ Allow SSH traffic from

Helps you connect to your instance

Anywhere

0.0.0.0/0

☒ Allow HTTPS traffic from the internet

To set up an endpoint, for example when creating a web server

☒ Allow HTTP traffic from the internet

To set up an endpoint, for example when creating a web server

8. Finally, press **Launch Instance** on the side.

The screenshot shows the AWS Management Console 'Launch Instance' wizard. The 'Configure storage' section is expanded, showing a single 8 GiB gp3 root volume. The 'Summary' section on the right shows 1 instance, Amazon Linux 2023 AMI, t2.micro instance type, and a new security group. The 'Launch Instance' button is highlighted in orange.

**Configure storage** [info](#) [Advanced](#)

1x 8 GiB gp3 Root volume (Not encrypted)

Free tier eligible customers can get up to 30 GB of EBS General Purpose (SSD) or Magnetic storage

Add new volume

Click refresh to view backup information  
The tags that you assign determine whether the instance will be backed up by any Data Lifecycle Manager policies.

0 x File systems [Edit](#)

**Advanced details** [info](#)

**Summary**

Number of instances [info](#)  
1

Software image (AMI)  
Amazon Linux 2023 AMI 2023.4.2...[read more](#)  
ami-09b90e09742640522

Virtual server type (instance type)  
t2.micro

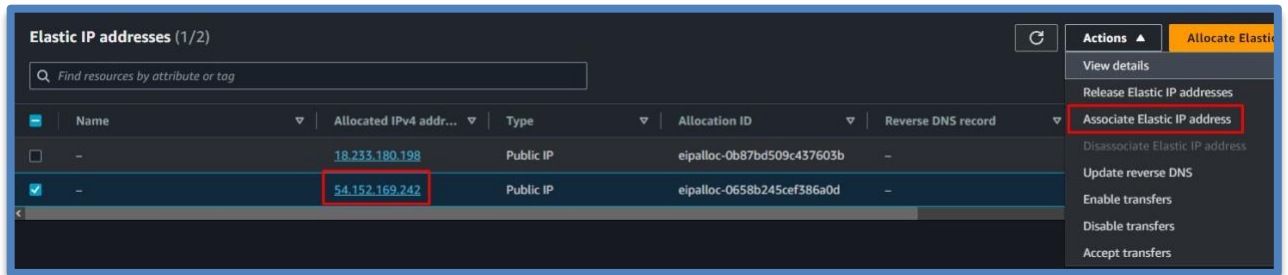
Firewall (security group)  
New security group

Storage (volumes)  
1 volume(s) - 8 GiB

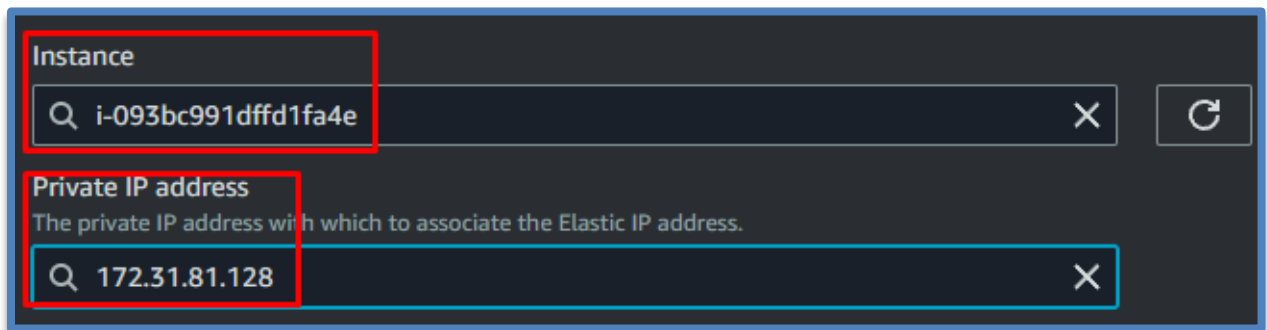
Free tier: In your first year includes 750 hours of t2.micro (or t3.micro in the Regions in which t2.micro is unavailable) instance usage on free tier AMIs per month, 750 hours of public IPv4 address usage per month, 30 GiB of EBS storage, 2 million IOs, 1 GB of snapshots, and 100 GB of bandwidth to the internet.

Cancel **Launch Instance** [Review commands](#)

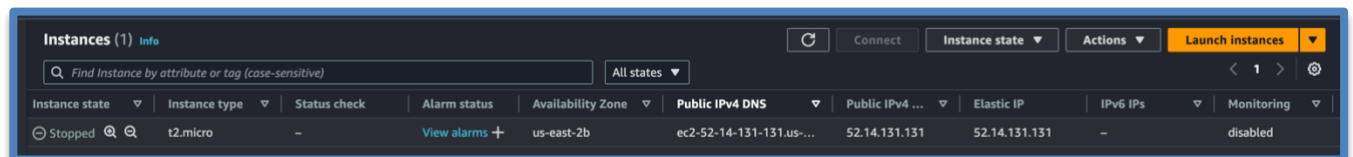
9. Before you connect to the instance, add the elastic IP to the instance. Go back to the Elastic IPs and select the one you got and press “Actions” and then “Associate Elastic IP Address”



10. Choose the instance that you just created and select the provided private IP, then press Associate.



11. When you go back to your instances, you should now see that the Elastic IP is now associated with the instance that you selected

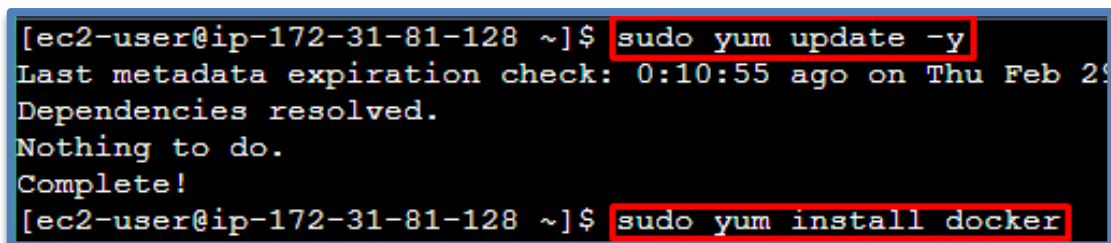


12. Connect to your instance with all the settings that are provided
13. Now that we are connected to the instance, we have to install docker on the instance to be able to run the commands. We use commands

```
sudo yum update -y
```

to make sure everything is up to date and then use

```
sudo yum install docker
```



14. Run the command
- ```
sudo service docker start
```
- To start the Docker service

15. Now run the command
- ```
sudo su
```

To get into the root account. You will run docker commands from this path

16. We now pull in our Docker image that we published using VS.

To do this, we run the command

```
docker pull $username/<project-name>:<tag>
```

in my case, this will be

```
docker pull ace1020302/scavengerus:dev
```

To make sure you got the correct one, you can run “docker images” and it will show all the images you have in the instance

```
[root@ip-172-31-81-128 ec2-user]# docker pull dmedwards2002/bucstop:latest
latest: Pulling from dmedwards2002/bucstop
5d0aeceef7ee: Pull complete
7c2bfda75264: Pull complete
950196e58fe3: Pull complete
ecf3c05ee2f6: Pull complete
819f3b5e3ba4: Pull complete
20ae32215d86: Pull complete
4f4fb700ef54: Pull complete
58b9c250b831: Pull complete
Digest: sha256:d9ee3fd5ac8ff0ee3686a99dccf5a84273c84df536737221480dadba855b35fc
Status: Downloaded newer image for dmedwards2002/bucstop:latest
docker.io/dmedwards2002/bucstop:latest
[root@ip-172-31-81-128 ec2-user]# docker images
```

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
dmedwards2002/bucstop	latest	5a80cb85901b	54 minutes ago	218MB

17. Next, we run the image using

```
docker run -d -p<port>:<port> $username/<project-name>:<tag>
```

for me it will look like

```
docker run -d -p80:80 ace1020302/scavengerus:dev
```

```
docker run -d -p80:80 dmedwards2002/bucstop:latest
```

18. (Optional) To make sure that the image is running run:

```
docker container ls
```

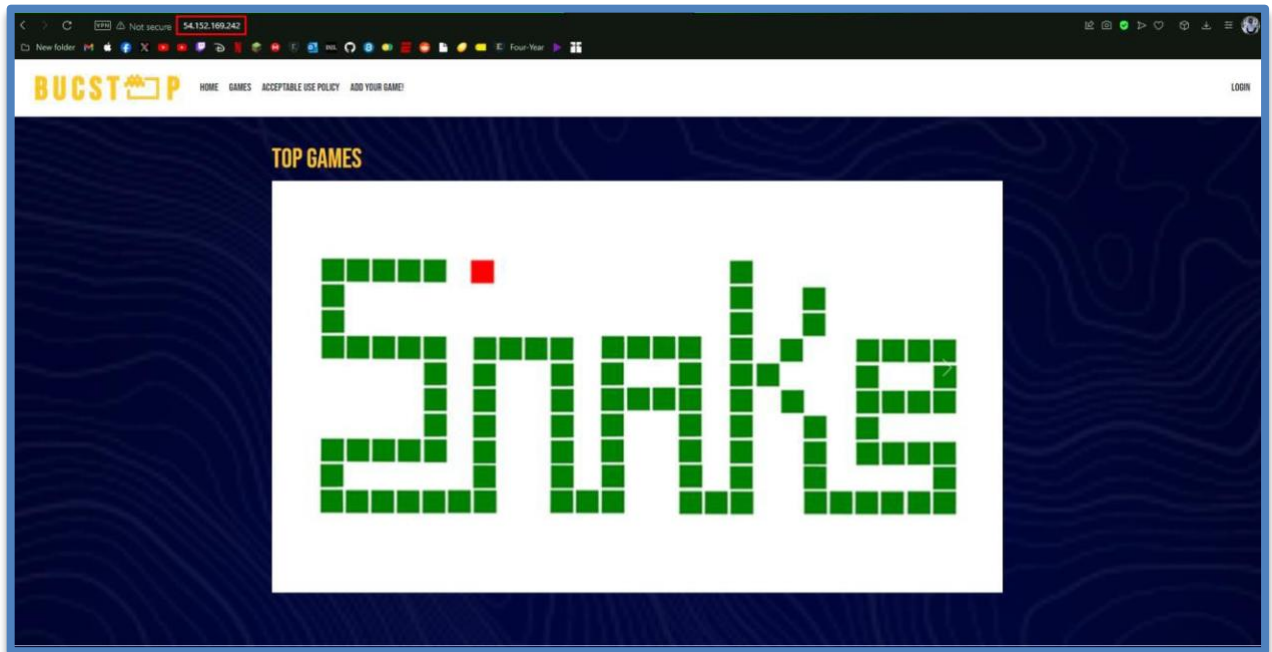
This will list all of your active containers

19. Now, to make sure it all worked, you can copy and paste that Elastic IP address into your address bar and press enter! If you do not remember what the IP is, then at the bottom of the instance, it will tell you what your public IP address is

i-093bc991dffd1fa4e (Test)

PublicIPs: 54.152.169.242 PrivateIPs: 172.31.81.128

20. And now you can give this IP to anyone while you have the container running, and they will be able to see your project!



Note: The container will continue running and that uses resources, resources that are limited due to AWS having a free version and if you go over the limits then you start paying. To stop this from happening, be sure to pause the container whenever you are done running it. You can also stop the container, but that gets rid of it entirely so pausing does the job just fine.

To do this, enter the command

```
docker container pause <container name>
```

for me it looks like

```
docker container pause exciting_lamport
```

AWS comes up with crazy names

```
[root@ip-172-31-81-128 ec2-user]# docker container ls
CONTAINER ID   IMAGE                                COMMAND                  CREATED        STATUS        PORTS                               NAMES
1a8608934226   dmedwards2002/bucstop:latest       "dotnet BucStop.dll"    8 minutes ago Up 8 minutes  0.0.0.0:80->80/tcp, :::80->80/tcp, 443/tcp exciting_lamport
[root@ip-172-31-81-128 ec2-user]# docker container pause exciting_lamport
exciting_lamport
[root@ip-172-31-81-128 ec2-user]# docker container ls
CONTAINER ID   IMAGE                                COMMAND                  CREATED        STATUS        PORTS                               NAMES
1a8608934226   dmedwards2002/bucstop:latest       "dotnet BucStop.dll"    8 minutes ago Up 8 minutes (Paused) 0.0.0.0:80->80/tcp, :::80->80/tcp, 443/tcp exciting_lamport
[root@ip-172-31-81-128 ec2-user]#
```

That is how you utilize AWS and Docker to deploy your project. Hope this helps you in figuring it all!